How nonconvex is it? Exploring neural network geometry through Dynamic String Sampling

C. Daniel Freeman[1, *] and Joan Bruna[2]

[1]*Department of Physics, University of California, Berkeley, CA 94720, USA*
[2]*Department of Statistics, University of California, Berkeley, CA 94720, USA*
(Dated: August 16, 2016)

The natural tool for understanding neural network training is *nonconvex optimization*. In practice, this means that, for some given learning task, a loss function is optimized via some flavor of gradient descent. While this strategy has led to great success, it is surprising that not much is known about the geometry of the loss function itself. In this manuscript, we examine the geometry of these loss functions for several learning tasks (QUAD, MNIST). In so doing, we provide both a quantitative handle on problem nonconvexity as well as an algorithm for estimating that nonconvexity (Dynamic String Sampling). Operationally, given two models which have nearly the same loss L_0 on test data, our algorithm searches for continuous paths in the space of model parameters which do not exceed L_0. Whether such a path can be found defines a notion of connectedness, and thus nonconvexity.

I. INTRODUCTION

II. QUANTIFYING NONCONVEXITY

A. Definitions

For a model with network parameters θ_i, and a learning problem with sample space X, the fundamental object of study is the loss function, $L(X, \theta_i)$. In practice, one only has access to an estimate of the loss function over some restricted subset, χ_i, of the sample space: $E(L(X, \theta_i), \chi_i)$. Unless otherwise stated, the loss functions computed throughout are assumed to be on a restricted test set not used during training.

A key ingredient of the algorithm is the use of an *interpolated model*. For two given models, θ_1 and θ_2, we defined the interpolated model with parameter t as follows:

$$\Theta(\theta_1, \theta_2, t) := \theta_1(1 - t) + \theta_2 t \qquad (1)$$

Thus, the interpolated model parameters—i.e., weights and biases—are simply linearly interpolated between two given models.

Additionally, the algorithm requires an estimate of the interpolated loss curve:

$$\gamma(\theta_1, \theta_2) := L(X, \Theta(\theta_1, \theta_2, t)), t \in [0, 1] \qquad (2)$$

or, an estimate of the loss on those models which are linear interpolations sitting between θ_1 and θ_2. More specifically, we seek efficient estimates of the location of the maxima, $t^* := \left.\frac{d\gamma(\theta_1, \theta_2, t)}{dt}\right|_{t^*} = 0, \left.\frac{d^2\gamma(\theta_1, \theta_2, t)}{dt^2}\right|_{t^*} < 0$. While in principle, the interpolated loss curve could have rich structure, in practice it is generally fairly smooth, thus straightforward hill climbing algorithms can efficiently locate these points.

Finally, for a pair of models (θ_1, θ_2), it will be convenient to define the maximum interpolated error:

$$\Gamma(\theta_1, \theta_2) := \min_{\Theta^*(\theta_1, \theta_2)} \max \left. L(X, \theta_i) \right|_{\theta_i \in \Theta^*} \qquad (3)$$

where $\Theta^*(\theta_1, \theta_2)$ is *any* continuous path in the space of weights connecting θ_1 and θ_2. Thus, $\Gamma(\theta_1, \theta_2)$ represents the minimum possible maximum loss achieved by those paths in the space of weights connecting θ_1 and θ_2. More intuitively, if $\Gamma(\theta_1, \theta_2) \leq \max(L(X, \theta_1), L(X, \theta_2))$, then the models are "connected"—there exists a continuous path in the space of models with total loss never exceeding the maximum loss achieved by θ_1 or θ_2.

B. The Greedy Algorithm

1. Train two models θ_i and θ_j to a threshold loss value, L_0.

2. Determine the location of the global maxima, t^*, on the interpolated loss curve $\gamma(\theta_i, \theta_j)$.

3. Perform gradient descent on the interpolated model $\Theta(\theta_i, \theta_j, t^*) := \theta_{i,j}$ until it is below αL_0 for some $\alpha \in [0, 1]$.

4. Calculate the maxima of the interpolated losses $\gamma(\theta_i, \theta_{i,j})$ and $\gamma(\theta_{i,j}, \theta_j)$. If these maxima are below L_0, then stop recursing on this branch and proceed to remaining branches(see 5). If not, proceed to step 5.

5. For those pairs, θ_a, θ_b from step 4 for which the maxima exceeds L_0, start a new branch by returning to step 2 and making the replacement $i->a$ and $j->b$. If depth exceeds d, stop (see below).

We provide a cartoon of the algorithm in Fig. 1. If the algorithm succeeds, then the output of the algorithm is a sequence of models, θ_i such that the pairwise interpolated loss curve between each in a sequence will be less than the threshold L_0. Thus, the algorithm outputs a continuous path in parameter space connecting the original two models such that everywhere along the path, the

FIG. 1. A cartoon of the algorithm. $a)$: The initial two models with approximately the same loss, L_0. $b)$: The interpolated loss curve, in red, and its global maximum, occuring at $t = t^*$. $c)$: The interpolated model $\Theta(\theta_i, \theta_j, t^*)$ is added and labeled $\theta_{i,j}$. $d)$: Stochastic gradient descent is performed on the interpolated model until its loss is below αL_0. $e)$: New interpolated loss curves are calculated between the models, pairwise on a chain. $f)$: As in step $c)$, a new model is inserted at the maxima of the interpolated loss curve between θ_i and $\theta_{i,j}$. $g)$: As in step $d)$, gradient descent is performed until the model has low enough loss.

total loss is less than or equal to the loss of the original models.

As written, if a path does *not* exist, then the algorithm will clearly not converge. Thus, on top of the parameter α, discussed below, the algorithm has an additional free parameter in the *depth* chosen to explore. For convenience, we define the string of models produced by the algorithm at depth d with parameter α to be the *interpolated string*, $S(\theta_1, \theta_2, \alpha, d)$. These are precisely those models recursively generated by the algorithm in step 3. Further, these models are naturally ordered along a path, starting from θ_1 and terminating on θ_2, as indicated in Fig. 1.

Finally, to use this as a tool to diagnose convexity, we define the *maximum interpolated error at depth d and tolerance α*:

$$\tilde{\Gamma}(\theta_1, \theta_2, d, \alpha) := \max \gamma(\theta_i, \theta_j, t) \qquad (4)$$
$$i, \ j \text{ neighbors in } S(\theta_1, \theta_2, \alpha, d)$$

where by "neighbors in $S(\theta_1, \theta_2, \alpha, d)$", we only mean that the models are immediately adjacent on the interpolating string. This quantity upper bounds the true maximum interpolated error, i.e. (3).

In summary: the algorithm recursively produces and trains new models lying on a continuous path in the space of model parameters, i.e. a string. Training via gradient

descent biases the path towards valleys on the loss surface, thus encouraging the loss along this path to be low. In practice, the parameter α is chosen to be less than 1 to aid convergence. We provide numerical and theoretical evidence for this choice in section SECTIONGOHERE.

C. Constrained Dynamic String Sampling

While the algorithm presented in Sec. II B is fast for sufficiently smooth families of loss surfaces with few saddle points, here we present a slightly modified version which, while slower, provides more control over the convergence of the string. Instead of training intermediate models via full SGD to a desired accuracy, intermediate models will be subject to a constraint that ensures they are "close" to the neighboring models on the string. Specifically, intermediate models will be constrained to the unique hyperplane in weightspace equidistant from its two neighbors. This is similar to a sort of "L_1 regularization" where the loss function for a given model on the string, θ_i, has an additional term $\tilde{L}(\theta) = L(\theta) + \zeta(\|\theta_{i-1} - \theta_i\| + \|\theta_{i+1} + \theta_i\|)$. The strength of the ζ regularization term controls the "springy-ness" of the weightstring. note: make this more precise, the hyperplane constraint is stronger than the L_1 constraint...L_1 only keeps the model in a ball close to the midpoint between the models.

Because adapting DSS to use this constraint is straightforward, here we will describe an alternative "breadth-first" approach wherein models are trained in parallel until convergence. This alternative approach has the advantage that it will indicate a disconnection between two models "sooner" insofar as it will be clear two models cannot be connected once the loss on either of the two initial models, θ_1 or θ_2, is less than $\Gamma(\theta_1, \theta_2)$. The precise geometry of the loss surface will dictate which approach to use in practice.

Given two random models σ_i and σ_j where $|\sigma_i - \sigma_j| < \kappa$, we aim to follow the evolution of the family of models connecting σ_i to σ_j. Intuitively, almost every continuous path in the space of random models connecting σ_i to σ_j has, on average, the same (high) loss. For simplicity, we choose to initialize the string to the linear segment interpolating between these two models. If this entire segment is evolved via gradient descent, the segment will either evolve into a string which is entirely contained in a basin of the loss surface, or some number of points will become fixed at a higher loss. These fixed points are difficult to detect directly, but will be indirectly detected by the persistence of a large interpolated loss between two adjacent models on the string.

The algorithm proceeds as follows:

(0.) Initialize model string to have two models, σ_i and σ_j.

1. Begin training all models to the desired loss, keeping the instantaneous loss of all models being trained approximately constant..

2. If the pairwise interpolated loss $\gamma(\sigma_n, \sigma_{n+1})$ exceeds a tolerance α_1, insert a new model at the maximum of the interpolated loss between these two models. For simplicity, this tolerance is chosen to be $(1 + \alpha_1^*)$ times the instantaneous loss of all other models on the string.

3. Repeat steps (1) and (2) until all models (and interpolated errors) are below a threshold loss L_0, or until a chosen failure condition (see II D).

D. Failure Conditions

While the algorithms presented will faithfully certify two models are connected if the algorithm converges, it is worth reemphasizing that they do not guarantee that two models are disconnected if the algorithm fails to converge. In general, the problem of determining if two models are connected can be made arbitrarily difficult by choice of a particularly pathological geometry for the loss function, so we are constrained to heuristic arguments for determining when to stop running the algorithm.

Thankfully, in practice, loss function geometries for problems of interest are not intractably difficult to explore.

III.   ANALYTIC TOY MODEL

To build intuition for the nonlinear case, here we treat an "almost linear" regression task. The strategy will be first, to understand how dynamic string sampling applies in an analytically tractable model, and second, to leverage this intuition for the more complicated numerical systems to follow. We will proceed by studying the level sets of the QUARTICLOSS toy model defined as followed:

For weight matrices w_1, w_2, input column vectors x_i of dimension d_1, and output scalars y_i, suppose we aim to minimize the loss function

$$L(w_1, w_2) = (w_1 w_2 x - y)^2 \tag{5}$$

For simplicity, suppose w_1 is a row vector with dimension d_2. This fixes the dimension of w_2 to d_1 by d_2.

Suppose we are given the globally optimal w_1^* and w_2^* which minimize L. These weights are not unique, because we can reparameterize the global optima in the following way:

$$w_1^* \rightarrow \text{any nonzero vector} = \text{w}_1 \tag{6}$$

$$w_2^* \rightarrow \frac{w_1^T}{||w_1||} a^* + \alpha (w_1^T)_\perp a^* \tag{7}$$

where a^* is the row vector $a^* = w_1^* w_2^*$ which globally optimizes (5), and α is an arbitrary scalar. This freedom represents an internal symmetry of the weights.

Given two models $w_1^A w_2^A = a^*$ and $w_1^B w_2^B = a^*$, let $w_1^A \cdot w_1^B = ||w_1^A|| ||w_1^B|| \cos\theta$. Either can be continuously

connected to the other via the following two steps. Without loss of generality, we will connect the A components to the B components. First, the component of w_2^A perpendicular to w_1 is scaled linearly from its original value α to α^*:

$$w_2^A = \frac{w_1^T}{||w_1||} a^* + (w_1^T)_\perp a^* \rightarrow \frac{w_1^T}{||w_1||} a^* + \alpha^* (w_1^T)_\perp =: w_2^{**} \tag{8}$$

Where $\alpha^* = \frac{1 - ||w_1^B|| \cos\theta}{||w_1^B|| \sin\theta}$. Defined in this way, w_2^{**} now satisfies both $w_1^A w_2^{**} = a^*$ as well as $w_1^B w_2^{**} = a^*$. Then, simple linear interpolation connects w_1^A to w_1^B:

$$w_1^a \rightarrow w_1^a (1 - t) + w_1^b t \tag{9}$$

Finally, the same procedure can connect w_2^B to w_2^A, continuously and invertibly, thus we are done.

Thus, linearly scaling first α to α^* and then scaling t from 0 to 1 connects w_1^A to w_1^B. Crucially, the entire procedure preserves the global minimum of the loss function.

Now that we have shown a continuous path exists for this internal symmetry, it remains to be shown that DSS will faithfully produce a continuous path. For DSS, the relevant interpolation equations are:

$$w_1(t) = w_1^A (1 - t) + w_1^B t \tag{10}$$

$$w_2(t) = w_2^A (1 - t) + w_2^B t \tag{11}$$

Examining $w_1(t)w_2(t)$, the first and last terms in the product yield $(1-t)^2 a^* + t^2 a^*$. The remaining cross terms can be rewritten as $\gamma t(t-1) a^*$ for a constant γ. Because we are promised that A and B are global minima in the space of weights, the loss is either constant along the path, or reaches a reaches a global maxima somewhere, which we call t^*. In the former case, we are done, because DSS will have provided a continuous path in the space of weights with nonincreasing loss.

In the latter case, consider the hyperplane splitting A and B. Because the loss function is continuous, and because we have demonstrated the existence of a continuous path in weight space connecting A to B, this path must puncture the hyperplane somewhere. We assume, but do not prove here, that constrained gradient descent on the hyperplane will find this new global minima. We comment on this assumption more below. Call this point π_1, and the associated weights $w_1^{\pi_1}$ and $w_2^{\pi_1}$.

Recursively splitting the weight string and performing gradient descent on the resulting constraining hyperplanes provides a sequence of points (reordered sequentially from A to B) $\pi_1, \pi_2, ..., \pi_n$ where $n = \Sigma_{i=0}^{d} 2^i$ for depth d recursive splits. By the continuity of the loss function, the maximum interpolated error along the path $A, \pi_1, ..., \pi_n, B$ can be bounded above to $L_0 + \epsilon$ for any $\epsilon > 0$ so long as d is made sufficiently large, where L_0 is the globally minimum loss achieved by a^*.

This bound can be made more precise...

we note that $\nabla L\Big|_{t^*}$ ...

IV.   NUMERICAL EXPERIMENTS

For our numerical experiments, we aimed to extract qualitative features of both small, toy networks, as well as of larger workhorse networks suitable for use on real world tasks (e.g. MNIST). At its core, the maximum interpolated error (i.e., (3)) is a measure of problem nonconvexity—or, more precisely, of the nonconvexity of the loss surface of a given architecture on a particular learning problem.

A.   Polynomial Regression

Polynomial function regression is a task for which small neural networks can achieve extremely high accuracy. For our numerical experiments, we studied a 1-4-4-1 fully connected multilayer perceptron style architecture with RELU activation and RMSProp optimization. For ease-of-analysis, we restricted the family of polynomials to be strictly contained in the interval $x \in [0,1]$ and $f(x) \in [0,1]$.

Discussion of different Loss functions
etc.

B.   MNIST

V.   DISCUSSION